

```
pip install ucimlrepo
```

```
Collecting ucimlrepo
  Downloading ucimlrepo-0.0.7-py3-none-any.whl.metadata (5.5 kB)
Requirement already satisfied: pandas>=1.0.0 in /usr/local/lib/python3.11/dist-packages (from ucimlrepo) (2.2.2)
Requirement already satisfied: certifi>=2020.12.5 in /usr/local/lib/python3.11/dist-packages (from pandas>=1.0.0->ucimlrepo) (2025.1.31)
Requirement already satisfied: numpy>=1.23.2 in /usr/local/lib/python3.11/dist-packages (from pandas>=1.0.0->ucimlrepo) (2.0.2)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas>=1.0.0->ucimlrepo) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas>=1.0.0->ucimlrepo) (2025.1)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas>=1.0.0->ucimlrepo) (2025.1)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-packages (from python-dateutil>=2.8.2->pandas>=1.0.0->ucimlrepo) (1.17.0)
Downloading ucimlrepo-0.0.7-py3-none-any.whl (8.0 kB)
Installing collected packages: ucimlrepo
Successfully installed ucimlrepo-0.0.7
```

```
#First two steps are to bring in the data from the referenced website
from ucimlrepo import fetch_ucirepo
```

```
national_health_and_nutrition_health_survey_2013_2014_nhanes_age_prediction_subset = fetch_ucirepo(id=887)
```

```
X = national_health_and_nutrition_health_survey_2013_2014_nhanes_age_prediction_subset.data.features
y = national_health_and_nutrition_health_survey_2013_2014_nhanes_age_prediction_subset.data.targets
```

```
print(national_health_and_nutrition_health_survey_2013_2014_nhanes_age_prediction_subset.metadata)
```

```
print(national_health_and_nutrition_health_survey_2013_2014_nhanes_age_prediction_subset.variables)
```

```
{'uci_id': 887, 'name': 'National Health and Nutrition Health Survey 2013-2014 (NHANES) Age Prediction Subset', 'repository_url': 'https://www.uci.edu/ml/repo0181/l01/nhanes_age_prediction_subset.zip'}
```

	name	role	type	demographic
0	SEQN	ID	Continuous	None
1	age_group	Target	Categorical	Age
2	RIDAGEYR	Other	Continuous	Age
3	RIAGENDR	Feature	Continuous	Gender
4	PAQ605	Feature	Continuous	None
5	BMXBMI	Feature	Continuous	None
6	LBXGLU	Feature	Continuous	None
7	DIQ010	Feature	Continuous	None
8	LBXGLT	Feature	Continuous	None
9	LBXIN	Feature	Continuous	None

  

	description	units	missing_values
0	Respondent Sequence Number	None	no
1	Respondent's Age Group (senior/non-senior)	None	no
2	Respondent's Age	None	no
3	Respondent's Gender	None	no
4	If the respondent engages in moderate or vigorous physical activity	None	no
5	Respondent's Body Mass Index	None	no
6	Respondent's Blood Glucose after fasting	None	no
7	If the Respondent is diabetic	None	no
8	Respondent's Oral Glucose Tolerance Test	None	no
9	Respondent's Blood Insulin Levels	None	no

```
import pandas as pd
import numpy as np
```

```
#Create a dataframe with Features and Targets
national = national_health_and_nutrition_health_survey_2013_2014_nhanes_age_prediction_subset
```

```
df = pd.DataFrame(national.data.features)
```

```
df = df.assign(age_group=national.data.targets)
```

```
df = df.assign(RIDAGEYR=national.data.targets)
```

```
#Index the insulin levels to create a new column with assigned labels based off the value in the insulin column
df['Category'] = pd.cut(df['LBXIN'],
                        bins=[0, 10, 15, 30, 110],
                        labels=['Low', 'Medium', 'High', 'Very High'],
                        right=False,
                        include_lowest=True)
```

```
df.head()
```

	RIAGENDR	PAQ605	BMXBMI	LBXGLU	DIQ010	LBXGLT	LBXIN	age_group	RIDAGEYR	Category
0	2.0	2.0	35.7	110.0	2.0	150.0	14.91	Adult	Adult	Medium
1	2.0	2.0	20.3	89.0	2.0	80.0	3.85	Adult	Adult	Low
2	1.0	2.0	23.2	89.0	2.0	68.0	6.14	Adult	Adult	Low
3	1.0	2.0	28.9	104.0	2.0	84.0	16.15	Adult	Adult	High
4	2.0	1.0	35.9	103.0	2.0	81.0	10.92	Adult	Adult	Medium

Next steps: [Generate code with df](#) [View recommended plots](#) [New interactive sheet](#)

```
#Using Label encoding to change Adult and Senior into 1 and 2 for the KNN
```

```
d = {'Adult': 1, 'Senior': 2}
df['age_group'] = df['age_group'].map(d)
df['RIDAGEYR'] = df['RIDAGEYR'].map(d)
print(df)
```

	RIAGENDR	PAQ605	BMXBMI	LBXGLU	DIQ010	LBXGLT	LBXIN	age_group	\
0	2.0	2.0	35.7	110.0	2.0	150.0	14.91	1	
1	2.0	2.0	20.3	89.0	2.0	80.0	3.85	1	
2	1.0	2.0	23.2	89.0	2.0	68.0	6.14	1	
3	1.0	2.0	28.9	104.0	2.0	84.0	16.15	1	
4	2.0	1.0	35.9	103.0	2.0	81.0	10.92	1	
...	...	...	...	...	...	...	...	...	
2273	2.0	2.0	33.5	100.0	2.0	73.0	6.53	1	
2274	1.0	2.0	30.0	93.0	2.0	208.0	13.02	1	
2275	1.0	2.0	23.7	103.0	2.0	124.0	21.41	1	
2276	2.0	2.0	27.4	90.0	2.0	108.0	4.99	1	
2277	1.0	2.0	24.5	108.0	2.0	108.0	3.76	1	

	RIDAGEYR	Category
0	1	Medium
1	1	Low
2	1	Low
3	1	High
4	1	Medium
...	...	...
2273	1	Low
2274	1	Medium
2275	1	High
2276	1	Low
2277	1	Low

[2278 rows x 10 columns]

```
import pandas
from sklearn import tree
from sklearn.tree import DecisionTreeClassifier
import matplotlib.pyplot as plt
```

```
#To give more familiar names to the column headers
df.rename(columns={'RIAGENDR': 'Gender', 'PAQ605': 'Exercise', 'BMXBMI': 'BMI', 'LBXGLU': 'Glucose', 'DIQ010': 'Diabetic', 'LBXGLT': 'Oral_G
```

```
print(df)
```

	Gender	Exercise	BMI	Glucose	Diabetic	Oral_Glucose	Insulin	\
0	2.0	2.0	35.7	110.0	2.0	150.0	14.91	
1	2.0	2.0	20.3	89.0	2.0	80.0	3.85	
2	1.0	2.0	23.2	89.0	2.0	68.0	6.14	
3	1.0	2.0	28.9	104.0	2.0	84.0	16.15	
4	2.0	1.0	35.9	103.0	2.0	81.0	10.92	
...	...	...	...	...	...	...	...	
2273	2.0	2.0	33.5	100.0	2.0	73.0	6.53	
2274	1.0	2.0	30.0	93.0	2.0	208.0	13.02	
2275	1.0	2.0	23.7	103.0	2.0	124.0	21.41	
2276	2.0	2.0	27.4	90.0	2.0	108.0	4.99	
2277	1.0	2.0	24.5	108.0	2.0	108.0	3.76	
	age_group	Age	Insulin_Category					
0	1	1	Medium					
1	1	1	Low					
2	1	1	Low					
3	1	1	High					

4	1	1	Medium
...	...	...	...
2273	1	1	Low
2274	1	1	Medium
2275	1	1	High
2276	1	1	Low
2277	1	1	Low

[2278 rows x 10 columns]

```
#Using Z-Score to detect outlier and then remove them from columns Glucose, BMI and Oral Glucose
from scipy import stats
```

```
import numpy as np
```

```
z = np.abs(stats.zscore(df['Glucose']))
```

```
print(z)
```

```
0      0.584085
1      0.590024
2      0.590024
3      0.248625
4      0.192715
...
2273    0.024985
2274    0.366384
2275    0.192715
2276    0.534114
2277    0.472265
Name: Glucose, Length: 2278, dtype: float64
```

```
threshold_z = 2
```

```
outlier_indices = np.where(z > threshold_z)[0]
```

```
no_outliers = df.drop(outlier_indices)
```

```
print("Original DataFrame Shape:", df.shape)
```

```
print("DataFrame Shape after Removing Outliers:", no_outliers.shape)
```

```
Original DataFrame Shape: (2278, 10)
DataFrame Shape after Removing Outliers: (2237, 10)
```

```
z = np.abs(stats.zscore(no_outliers['BMI']))
```

```
print(z)
```

```
0      1.092822
1      1.050235
2      0.646672
3      0.146537
4      1.120654
...
2273    0.786671
2274    0.299613
2275    0.577092
2276    0.062202
2277    0.465765
Name: BMI, Length: 2237, dtype: float64
```

```
threshold_z = 2
```

```
outlier_indices = no_outliers[z > threshold_z].index
```

```
no_outliers2 = no_outliers.drop(outlier_indices)
```

```
print("Original DataFrame Shape:", no_outliers.shape)
```

```
print("DataFrame Shape after Removing Outliers:", no_outliers2.shape)
```

```
Original DataFrame Shape: (2237, 10)
DataFrame Shape after Removing Outliers: (2132, 10)
```

```
z = np.abs(stats.zscore(no_outliers2['Oral_Glucose']))
```

```
print(z)
```

```
0      1.016329
1      0.825439
2      1.141170
3      0.720195
4      0.799128
...
2273    1.009616
2274    2.542365
2275    0.332244
2276    0.088732
```

```
2277    0.088732
Name: Oral_Glucose, Length: 2132, dtype: float64
```

```
threshold_z = 2
```

```
outlier_indices = no_outliers2[z > threshold_z].index
df = no_outliers2.drop(outlier_indices)
print("Original DataFrame Shape:", no_outliers2.shape)
print("DataFrame Shape after Removing Outliers:", df.shape)
```

```
Original DataFrame Shape: (2132, 10)
DataFrame Shape after Removing Outliers: (2032, 10)
```

```
print(df)
```

```
Gender  Exercise  BMI  Glucose  Diabetic  Oral_Glucose  Insulin \
0      2.0      2.0  35.7   110.0      2.0      150.0    14.91
1      2.0      2.0  20.3    89.0      2.0      80.0     3.85
2      1.0      2.0  23.2    89.0      2.0      68.0     6.14
3      1.0      2.0  28.9   104.0      2.0      84.0    16.15
4      2.0      1.0  35.9   103.0      2.0      81.0    10.92
...      ...      ...      ...      ...      ...      ...
2272    1.0      2.0  22.5    98.0      2.0      79.0     3.51
2273    2.0      2.0  33.5   100.0      2.0      73.0     6.53
2275    1.0      2.0  23.7   103.0      2.0     124.0    21.41
2276    2.0      2.0  27.4    90.0      2.0     108.0     4.99
2277    1.0      2.0  24.5   108.0      2.0     108.0     3.76
```

```
age_group  Age  Insulin_Category
0          1    1             Medium
1          1    1              Low
2          1    1              Low
3          1    1              High
4          1    1             Medium
...      ...  ...      ...
2272      1    1              Low
2273      1    1              Low
2275      1    1              High
2276      1    1              Low
2277      1    1              Low
```

```
[2032 rows x 10 columns]
```

```
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.datasets import load_iris
import numpy as np
import matplotlib.pyplot as plt
```

```
#Assign the features to the X axis without the target; assign the target to the Y
X = df.drop('Age', axis=1)
X = df.drop('Insulin_Category', axis=1)
Y = df['Insulin_Category']
```

```
#Train the test and train split
X_train, X_test, y_train, y_test = train_test_split(
    X, Y, test_size = 0.2, random_state=42)
```

```
#KNN accuracy score
knn = KNeighborsClassifier(n_neighbors=7)
knn.fit(X_train, y_train)
print(knn.score(X_test, y_test))
```

```
0.8157248157248157
```

```
#run the KNN and plot for the 8 features
neighbors = np.arange(1, 9)
train_accuracy = np.empty(len(neighbors))
test_accuracy = np.empty(len(neighbors))
```

```
for i, k in enumerate(neighbors):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)

    train_accuracy[i] = knn.score(X_train, y_train)
    test_accuracy[i] = knn.score(X_test, y_test)
```

```
test_accuracy[1] = min.score(x_test, y_test)

plt.plot(neighbors, test_accuracy, label = 'Testing dataset Accuracy')
plt.plot(neighbors, train_accuracy, label = 'Training dataset Accuracy')

plt.legend()
plt.xlabel('n_neighbors')
plt.ylabel('Accuracy')
plt.show()
```

